

Final Project Report of C/C++ Programming MA 251

He LIU

January 6, 2016

Instructor: Professor Li

1 Introduction

In this project, we are supposed to write a program which can be used to sort massive amounts of data. Since the size of data may be too large to load into RAM, it is not feasible to sort directly. Generally, there are two methods to do this.

One is creating a memory-mapped file. A memory-mapped file contains the contents of a file in virtual memory. This mapping between a file and memory space enables an application, including multiple processes, to modify the file by reading and writing directly to the memory. With the memory-mapped files that are associated with a source file on a disk, when the last process has finished working with the file, the data are saved to the source file on the disk.

Another is using external sorting, which is an application of divide and conquer strategy. One example of external sorting is the external merge sort algorithm, which sorts chunks that each fit in RAM, then merges the sorted chunks together. The merge algorithm only makes one pass sequentially through each of the chunks, each chunk does not have to be loaded completely.

In this project, I tried to use the first method in the beginning, but failed. Most of the sources programmed in C are based on Linux/UNIX, and I do not have enough time nor enough ability to migrate them into Win32. Then I turned to try the second method and got right result of sorting. This report is based on the second method, external sorting.

2 Program Logic

I made a flow diagram (Figure 1) to show the procedure of my program.

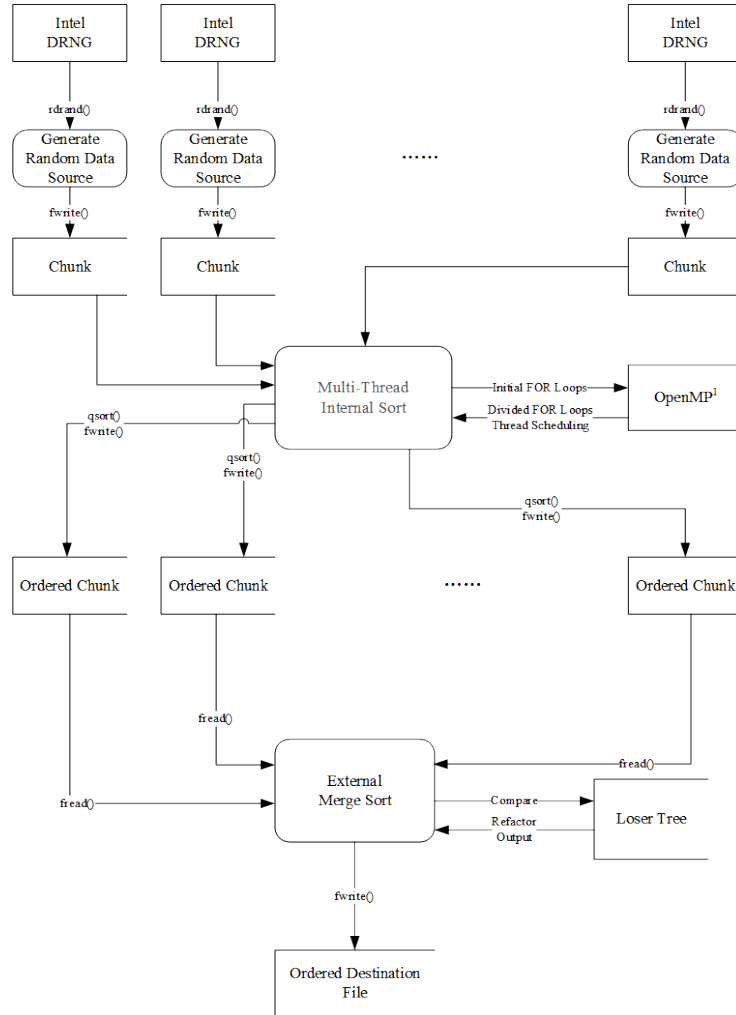


Figure 1: The procedure of my program.

¹OpenMP (Open Multi-Processing) is an application programming interface (API) that supports multi-platform shared memory multiprocessing.

3 Time Required

The integers involved in sorting are 32-bit integer ², which have a range of $-2\,147\,483\,648$ to $2\,147\,483\,647$.

The shortest time used in each case is shown in Table 1. ³

| Table 1: Shortest time required | | |
|---------------------------------|----------------|-------------------|
| NUM. OF INT32 | NUM. OF CHUNKS | TIME REQUIRED (s) |
| 1,024,000 | 8 | 0.225 |
| 5,120,000 | 8 | 0.673 |
| 10,240,000 | 8 | 1.216 |
| 51,200,000 | 16 | 5.746 |
| 102,400,000 | 32 | 11.623 |
| 204,800,000 | 64 | 23.162 |

All of the sorting results are checked by viewing them in Notepad++ with HEX-Editor plug-in.

4 Analysis

In this project, running time can be affected by internal sort (Quick Sort), file I/O and external sort. The function `qsort()` included in C STL has been heavily optimized and we do not need to do anything on this. However, there are many files to be sorted, if they can be sorted at the same time, lots of time can be saved. File I/O is a waste of time but cannot be removed, so the times to read or write files should be reduced as much as possible. The algorithm applied in external sort is important. In this project, we use Loser Tree to improve merge efficiency.

4.1 Parallel Quick Sort

There are 8 threads in the CPU of my computer. The default is each program only establish one process, but we can create multiple processes through programming. My basic idea on this is distributing the files waiting for being sorted in to 8 parts in average. Then let them do Quick Sort at the same time. Since thread scheduling takes some time, the speed up cannot be 8 times, but there is still much. Meanwhile, the time used on File I/O can be reduced.

In parallel coding, to reduce process latency, get the correct results and avoid

²The random integers are generated by `rand()`, which is included in Intel Digital Random Number Generator (DRNG), can generate “real” random integers based on CPUID . The initial random integers are 32-bit unsigned integers, which have a range of 0 to 4 294 967 295. I used `(u32 >> 16) & 0x7FFF` to convert them to 32-bit signed integers.

³The development environment is shown in Appendix A, and all of the raw data are attached in Appendix B.

variables competition, the design of parallel structure and the settings of variables are important.

In addition, this parallel computing follows forkjoin model, as show in Figure 2.




Figure 2: Fork-join Model

4.2 Loser Tree

In K-way merge, we can use Heap Sort or Winner/Loser Tree. When the number of ways is not too large, we often use Winner/Loser Tree to reduce the time used in adjust. Loser Tree and Winner Tree can find the extremum in $O(\log K)$, and the algorithm complexity of Loser Tree is $O((n - 1) \cdot \log K)$.

Loser Tree is an improvement from Winner Tree. Updating a Loser Tree is more convenient. The steps of Loser Tree algorithm is:

When a number is output, the number following enters as a leaf, and moves up by levels toward the root. As you move up, you carry the winner along. At each level, you find the loser of the last contest at that level. The node you are carrying challenges this previous loser. If it wins, you take it on to the next level. If it loses, you leave it and take the winner (the number which was there before you came) along to the next level. The winner of the contest at the root is placed the "extra" overall winner slot.

In addition, we can improve it by doing things in parallel: we can do the reading or writing at the same time as the sorting.

4.3 The link between the number of chunks and speed

As we can see in Appendix B, when the size of chunks increase, the speed may not be increased. And the reason is File I/O will consume more time. When the amount of data increase, more chunks will be needed because of the heavy

pressure on Quick Sort. If the size of source file is too large to load into RAM, it must be splitted into several chunks. I think there exists a appropriate number range of chunks to be the best one, but I do not have enough time to generate data and make curve fitting.

Since my computer has a 8 threads CPU, the minimal appropriate number should be no less than 8 in this “big data sorting”, and the number of chunks is supposed to be multiples of 8.

4.4 Others

- Use Binary File I/O.
- Do not use stack as a buffer for its size limit.
- Use QueryPerformanceCounter() in Win32API to print out the time used in a section.

5 Conclusion

In this project, I am more familiar with C programming, and learned a lot about sort algorithm. In addition, this project stimulated my interest. I tried to use parallel computing, tried to generate and sort 32-bit “real” random integers on Windows platform, tried to use Intel(R) Advisor and Intel(R) VTune(TM) Amplifier to help me optimize my program. I will try to use memory-mapped file on Linux platform later. I think the most valuable treasure I gained in this project is the interest on C programming.

References

- [1] <http://www.codingunit.com/writing-memory-to-a-file-and-reading-memory-from-a-file-in-c>
- [2] <http://chenkegarfield.blog.163.com/blog/static/62330008200910249526638/>
- [3] <https://software.intel.com/en-us/articles/the-drng-library-and-manual>

Appendices

A Development Environment

Table 2: Hardware

| Component | Model |
|-----------|---------------------------------|
| CPU | Intel(R) Core(TM) i7 4700HQ |
| SSD | TOSHIBA(R) Q200 EX. |
| RAM | HyperX(R) DDR3 1600MHz CL9 12GB |

Table 3: Software

| Software | Name |
|------------------------------|---|
| OS | Windows 10 TH2 x64 |
| IDE | Microsoft(R) Visual Studio 2015 Community |
| IDE Compile Mode | Release X86 |
| Compiler | Intel C++ Compiler 16.0 |
| Parallel Coding API | OpenMP 4.0 |
| Assisted optimization tool | Intel(R) Advisor XE 2016 |
| Assisted analysis tool | Intel(R) VTune™ Amplifier 2016 |
| Random integer generator API | Intel(R) Digital Random Number Generator |

B Raw Time Using Data

Table 4: Number of intergers: 1,024,000 Size of source file(MB): 4

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|--------------|
| 4 | 97 | 147 | 0.244 | 0.231 |
| | 71 | 152 | 0.223 | |
| | 74 | 151 | 0.225 | |
| 8 | 86 | 154 | 0.240 | 0.225 |
| | 87 | 137 | 0.224 | |
| | 61 | 151 | 0.212 | |
| 16 | 115 | 138 | 0.253 | 0.250 |
| | 100 | 145 | 0.245 | |
| | 109 | 143 | 0.252 | |
| 32 | 163 | 140 | 0.303 | 0.324 |
| | 172 | 157 | 0.329 | |
| | 186 | 154 | 0.340 | |
| 64 | 277 | 156 | 0.433 | 0.447 |
| | 252 | 158 | 0.410 | |
| | 343 | 155 | 0.498 | |

Table 5: Number of intergers: 5,120,000 Size of source file(MB): 20

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|--------------|
| 4 | 238 | 448 | 0.686 | 0.703 |
| | 248 | 463 | 0.711 | |
| | 244 | 468 | 0.712 | |
| 8 | 176 | 488 | 0.664 | 0.673 |
| | 177 | 492 | 0.669 | |
| | 177 | 508 | 0.685 | |
| 16 | 220 | 501 | 0.721 | 0.731 |
| | 232 | 483 | 0.715 | |
| | 242 | 515 | 0.757 | |
| 32 | 280 | 505 | 0.785 | 0.799 |
| | 309 | 507 | 0.816 | |
| | 281 | 514 | 0.795 | |
| 64 | 367 | 524 | 0.891 | 0.909 |
| | 409 | 506 | 0.915 | |
| | 398 | 524 | 0.922 | |

Table 6: Number of intergers: 10,240,000 Size of source file(MB): 40

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|--------------|
| 8 | 316 | 935 | 1.251 | 1.216 |
| | 304 | 900 | 1.204 | |
| | 299 | 893 | 1.192 | |
| 16 | 381 | 921 | 1.302 | 1.275 |
| | 329 | 920 | 1.249 | |
| | 325 | 948 | 1.273 | |
| 32 | 427 | 920 | 1.347 | 1.336 |
| | 397 | 930 | 1.327 | |
| | 383 | 952 | 1.335 | |
| 64 | 477 | 945 | 1.422 | 1.426 |
| | 467 | 961 | 1.428 | |
| | 495 | 934 | 1.429 | |

Table 7: Number of intergers: 51,200,000 Size of source file(MB): 200

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|--------------|
| 8 | 1681 | 4262 | 5.943 | 5.940 |
| | 1595 | 4295 | 5.890 | |
| | 1697 | 4289 | 5.986 | |
| 16 | 1381 | 4310 | 5.691 | 5.746 |
| | 1383 | 4472 | 5.855 | |
| | 1374 | 4317 | 5.691 | |
| 32 | 1445 | 4421 | 5.866 | 5.870 |
| | 1428 | 4454 | 5.882 | |
| | 1425 | 4436 | 5.861 | |
| 64 | 1559 | 4470 | 6.029 | 6.022 |
| | 1545 | 4471 | 6.016 | |
| | 1542 | 4479 | 6.021 | |

Table 8: Number of intergers: 102,400,000 Size of source file(MB): 400

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|---------------|
| 8 | 6236 | 8326 | 14.562 | 13.641 |
| | 5050 | 8478 | 13.528 | |
| | 4423 | 8411 | 12.834 | |
| 16 | 3078 | 8632 | 11.710 | 12.039 |
| | 3489 | 8762 | 12.251 | |
| | 3264 | 8892 | 12.156 | |
| 32 | 2805 | 8827 | 11.632 | 11.623 |
| | 2798 | 8946 | 11.744 | |
| | 2705 | 8789 | 11.494 | |
| 64 | 2820 | 8943 | 11.763 | 11.636 |
| | 2791 | 8790 | 11.581 | |
| | 2798 | 8767 | 11.565 | |
| 128 | 3178 | 9214 | 12.392 | 12.380 |
| | 3109 | 9156 | 12.265 | |
| | 3158 | 9325 | 12.483 | |

Table 9: Number of intergers: 204,800,000 Size of source file(MB): 800

| Num. of Chunks | Time of internal sort (ms) | Time of merge sort (ms) | Total time(s) | Avg. time(s) |
|-------------------|----------------------------------|-------------------------------|---------------|---------------|
| 8 | 8295 | 16680 | 24.975 | 24.985 |
| | 8237 | 16746 | 24.983 | |
| | 8220 | 16777 | 24.997 | |
| 16 | 7521 | 17241 | 24.762 | 25.124 |
| | 8596 | 17301 | 25.897 | |
| | 7617 | 17095 | 24.712 | |
| 32 | 7990 | 17687 | 25.677 | 25.912 |
| | 8797 | 17489 | 26.286 | |
| | 8357 | 17415 | 25.772 | |
| 64 | 5641 | 17573 | 23.214 | 23.162 |
| | 5555 | 17606 | 23.161 | |
| | 5577 | 17533 | 23.110 | |
| 128 | 5693 | 17958 | 23.651 | 23.809 |
| | 5704 | 18077 | 23.781 | |
| | 5662 | 18332 | 23.994 | |
| 256 | 6238 | 19218 | 25.456 | 26.122 |
| | 6172 | 18823 | 24.995 | |
| | 6237 | 21677 | 27.914 | |